

We didn't build a chatbot. We built a bank.

LATAM Bank keeps its word to every customer, by design.

-  Five promises, enforced in code
-  Run like a bank: owners, debate, audit
-  Live today, in Spanish and Portuguese

**Lía**
AI assistant, always disclosed

No reconozco este cargo

Encontré el cargo: 152.000 COP en Super Norte, el 2 de octubre.

Abrir reclamo
152.000 COP · tarjeta terminada en 0001

Confirmar Cancelar

 Hablar con una persona

The promise banks break every day

TODAY

37 hours

for a first answer after a customer reports a charge they never made



WITH LATAM BANK

Seconds

Lía finds the charge and proposes the claim. 95% of answers in under 5 seconds, measured offline

#1 complaint

Unrecognized charges: 12,297 of 67,095 complaints


3 in 4

of those complaints are still open

Minutes

is how fast fraud moves. Trust is lost in hours

We ran it like a bank: owners, debate, a record



21

research studies
before writing code

34

decisions, each with
its reason

10

architecture decision
records

66

sources tracked by
Audit

Governance challenged the design. Audit reviewed it cold. Every vice presidency owns its folder, its rules and its tests.

The disagreements are in the repository, not hidden behind the demo.

Five promises a bank can actually keep

1



**Nothing without
your yes**

Every action waits for
your approval on screen

2



Only your data

Other customers do not
exist for your session

3



**A person,
whenever you ask**

With your whole story.
You never repeat it

4



**Answers with a
source**

Policy cited by rule, or an
honest "I don't know"

5




**Only what was
verified**

It reports what the
system actually did



Enforced in code and versioned policy, not in a prompt. A prompt injection cannot break them.

Under the hood: the model talks, policy as code acts



MACHINE LEARNING, EACH VS A BASELINE



Policy retrieval with citation

0.90 vs 0.52 (BM25). In the workflow



Contact-reason classifier

macro-F1 0.40 vs 0.34 (rules)



SLA-breach risk

No signal (AUC 0.50). Reported, abstains

OPERATIONS



Tracing without message content



Least privilege, Secret Manager




Uptime check and alerts, Terraform



588 unit + 81 browser tests in CI

We tried to break it. Then we published the results



89%

of 129 runs pass. Rules engine: 72%

0

actions without approval, in 225 runs

0 / 18

missed handoffs to a person

1

real flaw in 129 runs. Found and published

**Customers: 37 hours to seconds**

**Experts: the full picture, drafted**

**The bank: auditable policy**